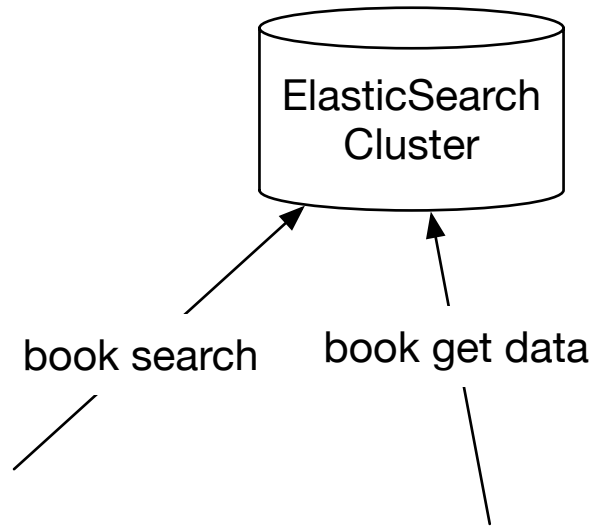



the sub-DBs keep copies of centralized repo.
There's two DBs for each library.
One, ES cluster of book metadata, driving search.
Two, a SQL DB keeping nothing but business data (i.e. loans, etc.).

Centralized repo is indexed by ISBN.
Only responds to ISBN get's.

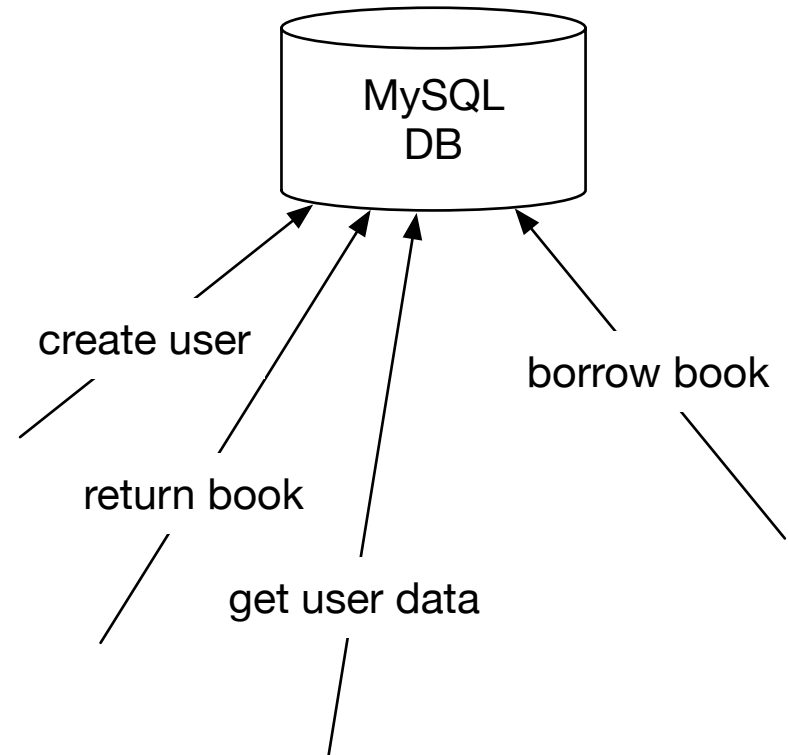
For handling search, metadata



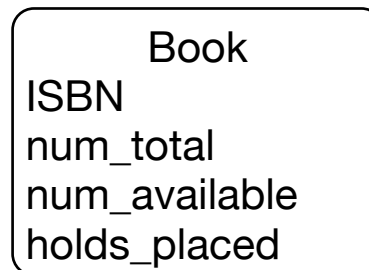
example of ES book



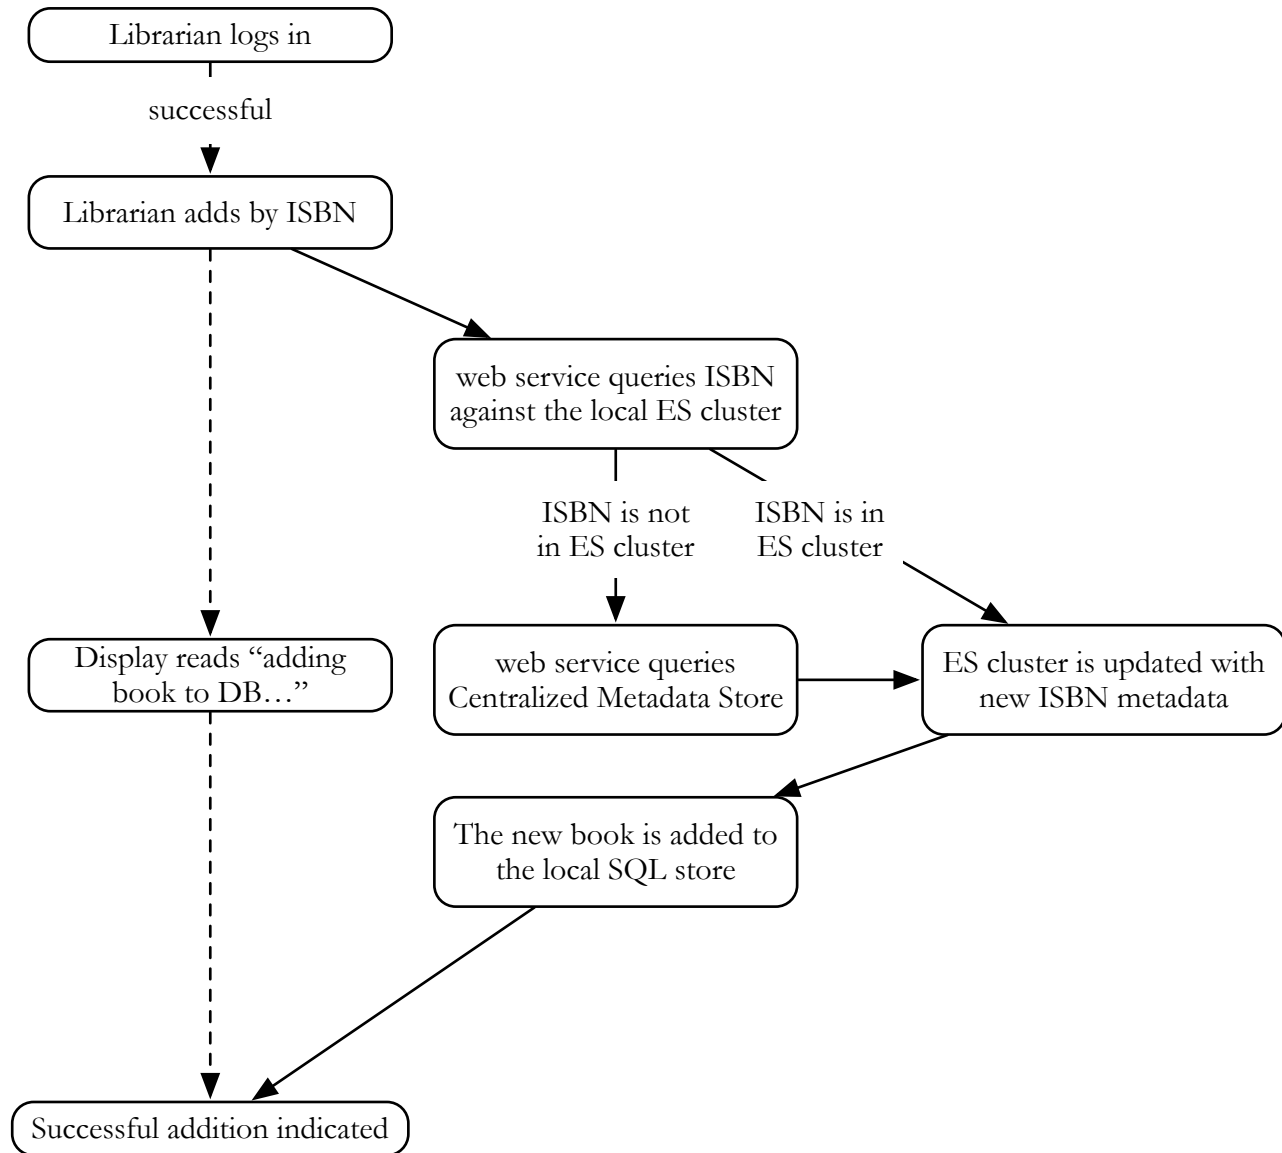
For handling loans, member data



e.g. of MySQL book



soul-searching question:
is it worthwhile to keep
separate records for
different copies of one book?



GETish Topics (Queries)

lid/query-book
lid/query-book-loan-status
lid/query-member-data

POSTish Topics (Modification Commands)

lid/loan-book
lid/return-book

PUTish Topics (Creation/Total Replacement Commands)

lid/add-book
lid/remove-book
lid/add-user
lid/remove-user

